

Process Abstraction

2023-24 COMP3230A

Contents

- ◉ What is a “process”?
- ◉ How to represent a “process”?
 - ◉ Resources use by a process
 - ◉ Process states
- ◉ Important data structures
- ◉ Operations on processes

Learning Outcome

- ◉ ILO 2a - **explain** how OS manages **processes**/threads and discuss the mechanisms and policies in efficiently sharing of CPU resources.

Reading & Reference

- ◉ Required Reading

- ◉ Operating Systems: Three Easy Pieces by Arpaci-Dusseau et. al
 - ◉ Chapter 4, Abstraction: The Process
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-intro.pdf>
 - ◉ Chapter 5, Interlude: Process API
 - ◉ <http://pages.cs.wisc.edu/~remzi/OSTEP/cpu-api.pdf>

- ◉ Reference

- ◉ Chapter 3 of Operating Systems, 3rd edition by Deitel et. al

Process vs. Program

- ◉ Program itself is a lifeless thing
- ◉ What is a Process?
 - ◉ A process is a program in execution
 - ◉ **An instance** of a program running in a computer
 - ◉ A process is an entity that can be assigned to and executed on a CPU
 - ◉ A unit of activity characterized by
 - ◉ the execution of a sequence of instructions,
 - ◉ with an execution state, and
 - ◉ an associated set of system resources

Process – an abstraction

- ◉ To represent a running program, the OS needs to keep track of the following information:

The memory that the process can access (or reference) is part of the process

Memory

Examples

Program code

Data

Resources
in use

I/O in use

Physical memory

During execution, process updates/ stores data in CPU registers, e.g., program counter, stack pointer, etc.

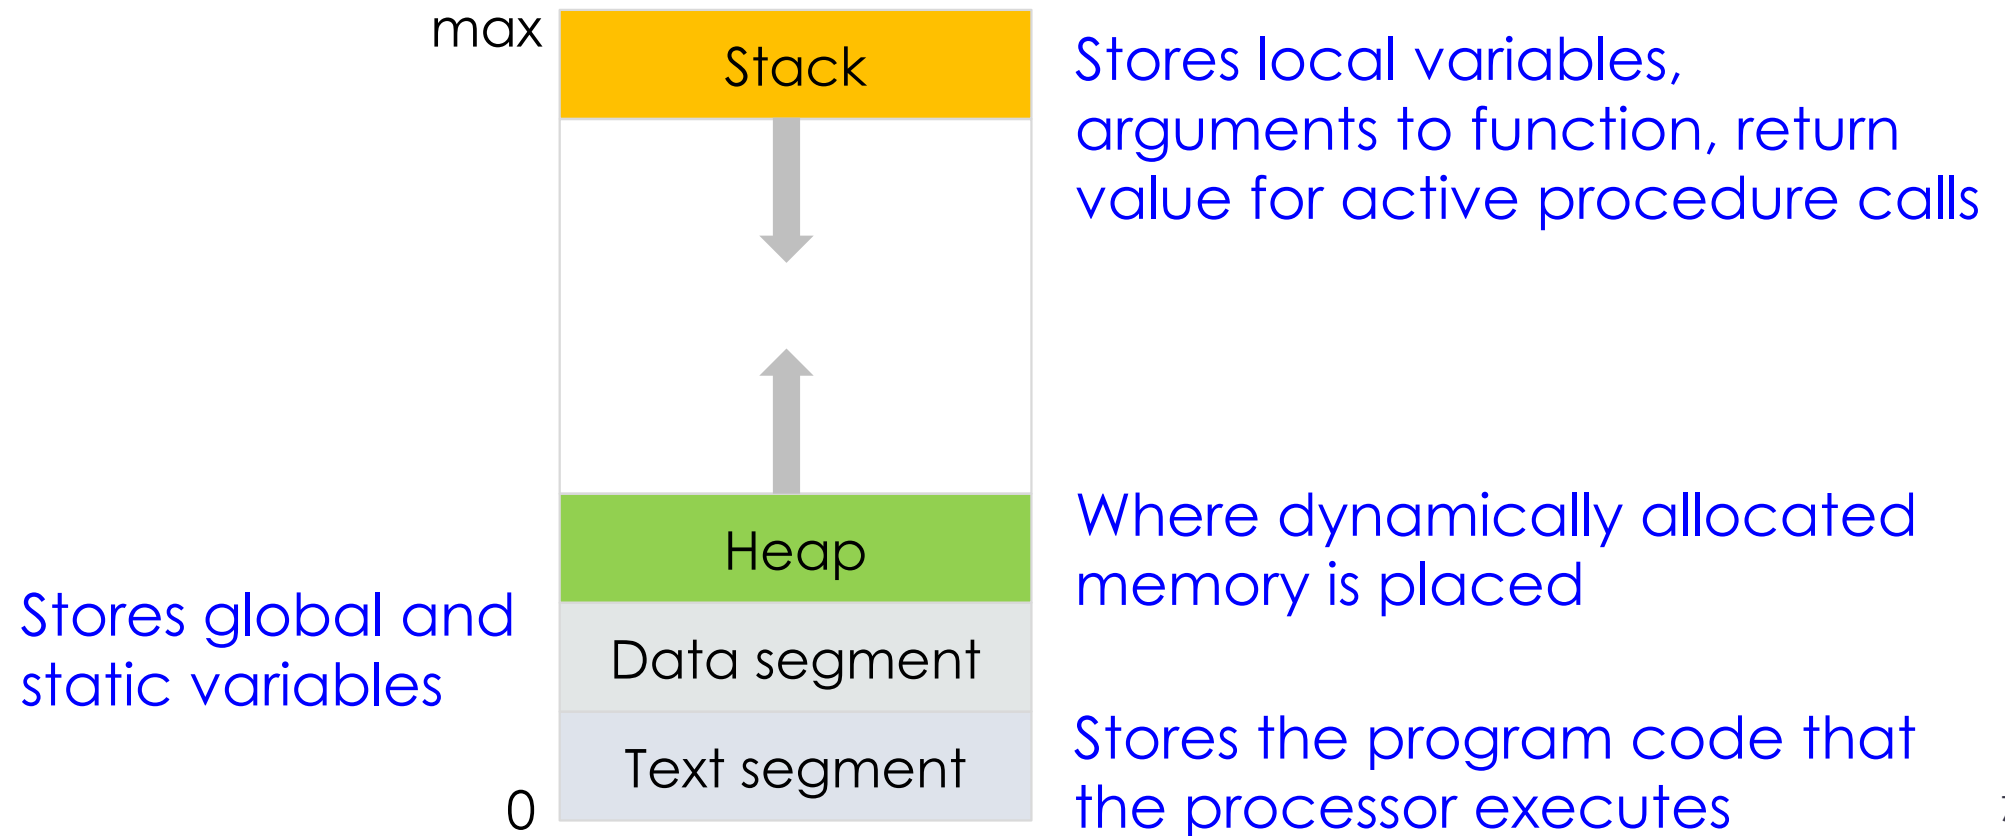
Execution
state

Register set

Current process
state

Process – Address Space

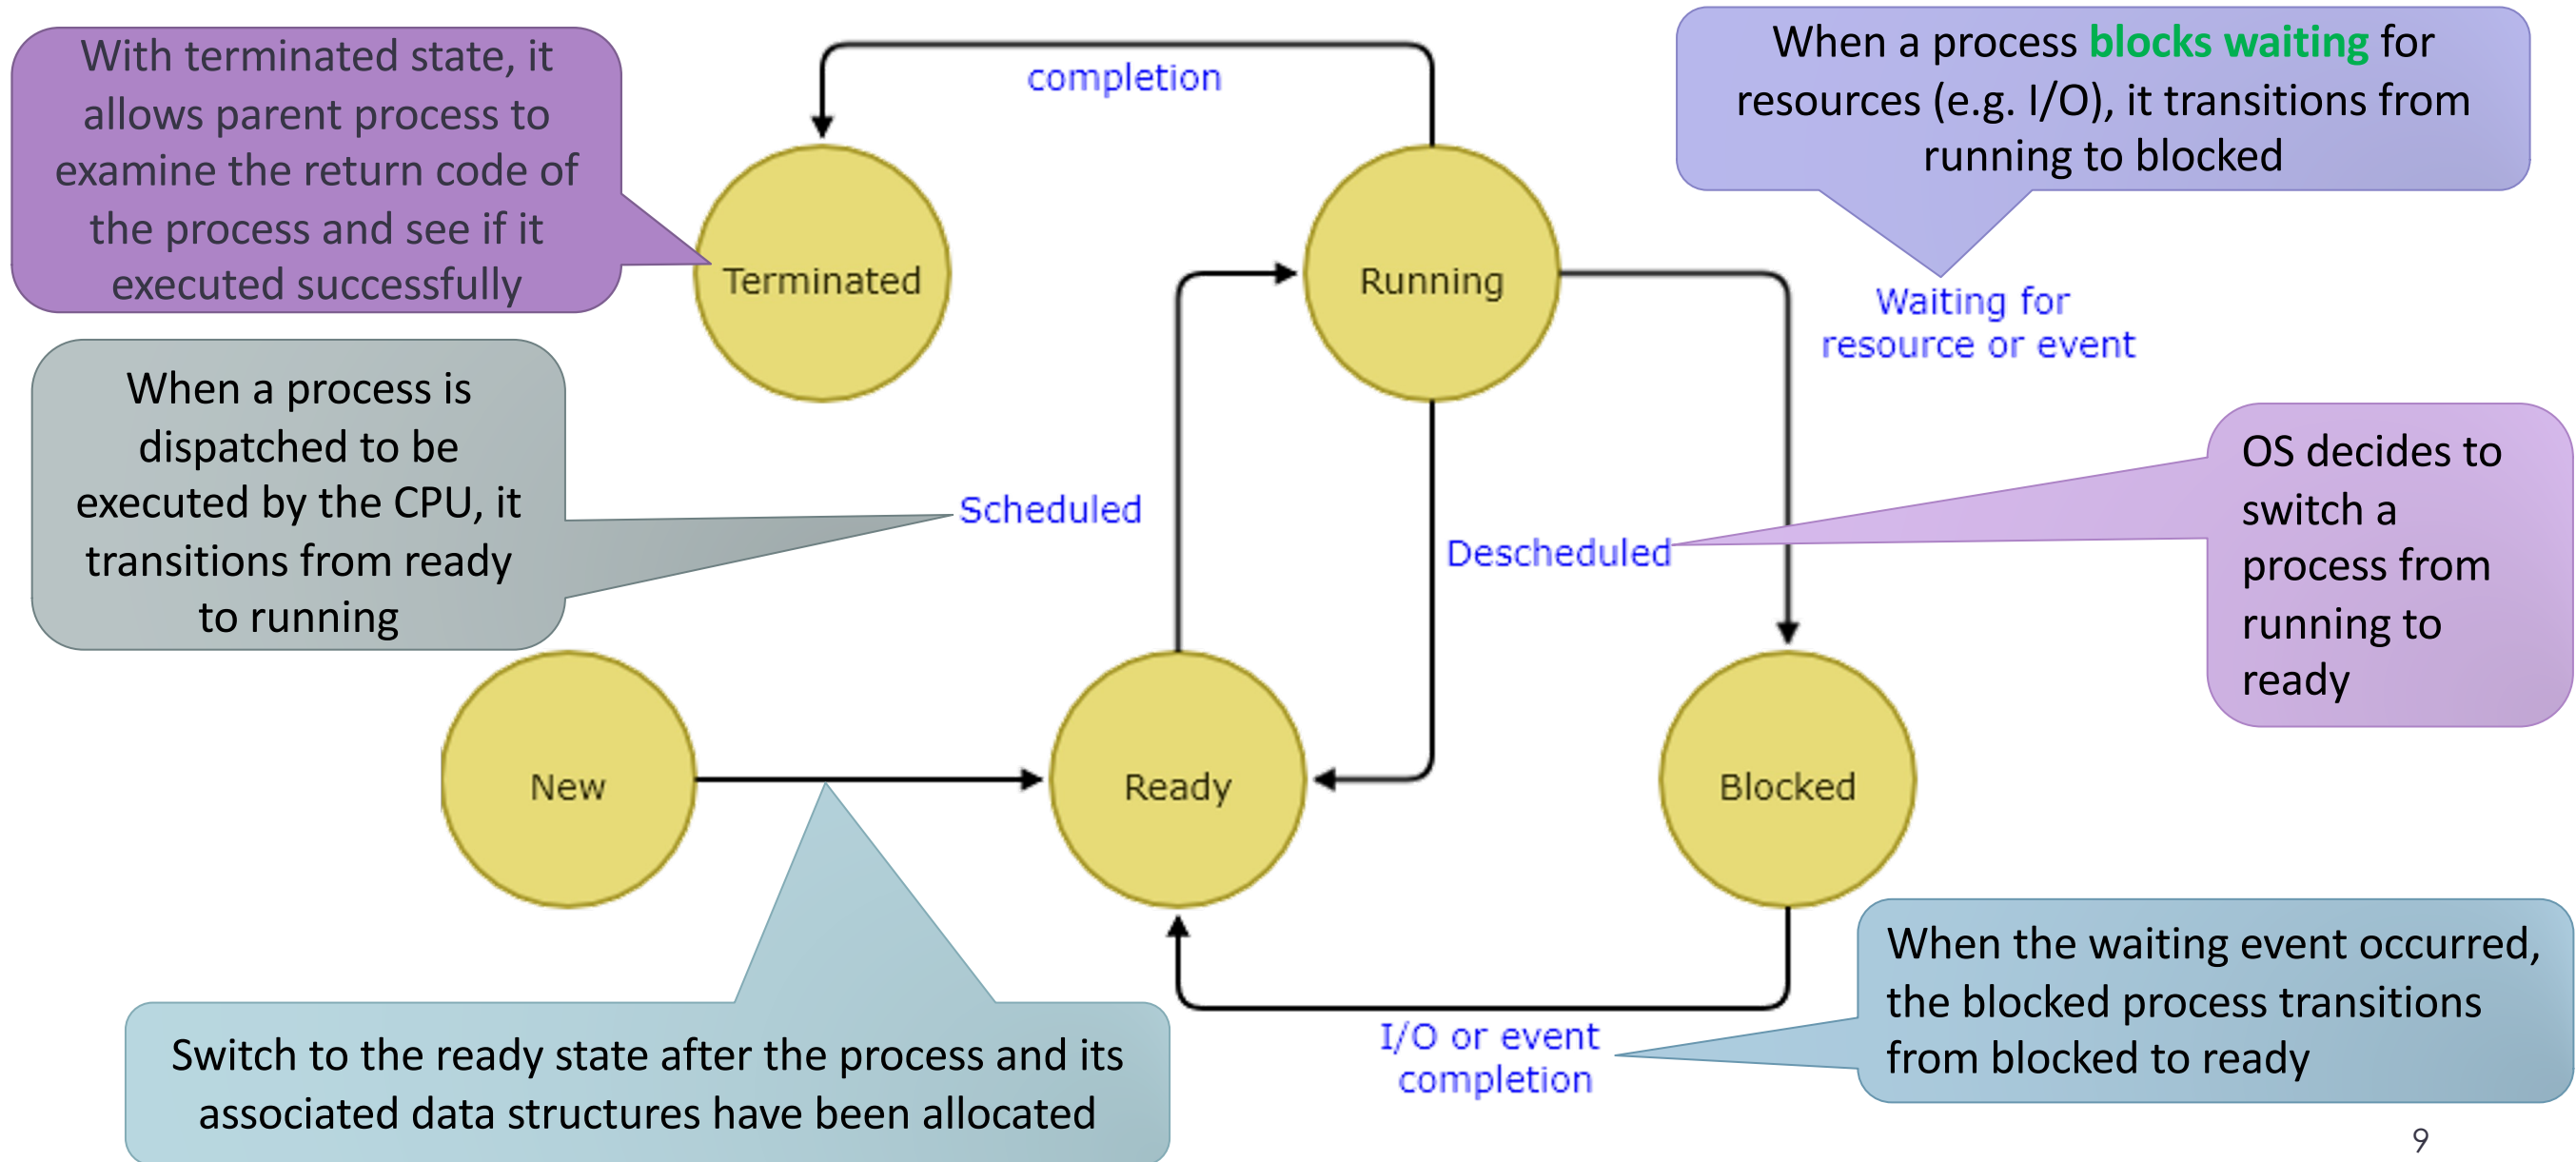
- ◉ The process's view of its memory is called the **address space**, which is a range of memory locations (or a range of memory addresses)
- ◉ Process address space consisting of a few “regions”:



Process – Process States

- ◉ In the process execution life cycle, it moves through a series of discrete process states.
- ◉ Process state - an indicator of the nature of the current activity of a process
 - ◉ **new (initial)**: The process is just being created
 - ◉ **running**: The process is executing on a processor
 - ◉ **blocked**: The process is **waiting for** some event (e.g., I/O or communication) to happen before it can proceed
 - ◉ **ready**: The process **is ready to run** on a processor and **is waiting** to be assigned to a processor
 - ◉ **terminated (final / zombie)**: The process has finished execution but has not yet been cleaned up; why not just discard it?

Life Cycle of a Process



Process Control Block

- ◉ To manage a process, OS makes use of a **data structure** to maintain information about a process
 - ◉ Process Control Block (**PCB**) or Process Descriptor
- ◉ PCB typically includes
 - ◉ Process identification number (PID) - **a unique ID**
 - ◉ Current **process state**
 - ◉ Program counter - indicates the address of **next** instruction
 - ◉ **Register context** - a snapshot of the register contents in which the process **was last running before** it **transitioned out of the running state**
 - ◉ Scheduling information - process priority, pointers to scheduling queues, etc.

Process Control Block (2)

- ◉ Credentials - determines the resources this process can access
- ◉ Memory Management information - concerning memory areas allocated to the process
- ◉ Accounting information - CPU usage statistics, time limits, etc.
- ◉ A pointer to the process's parent process
- ◉ Pointers to the process's child processes
- ◉ Pointers to allocated resources

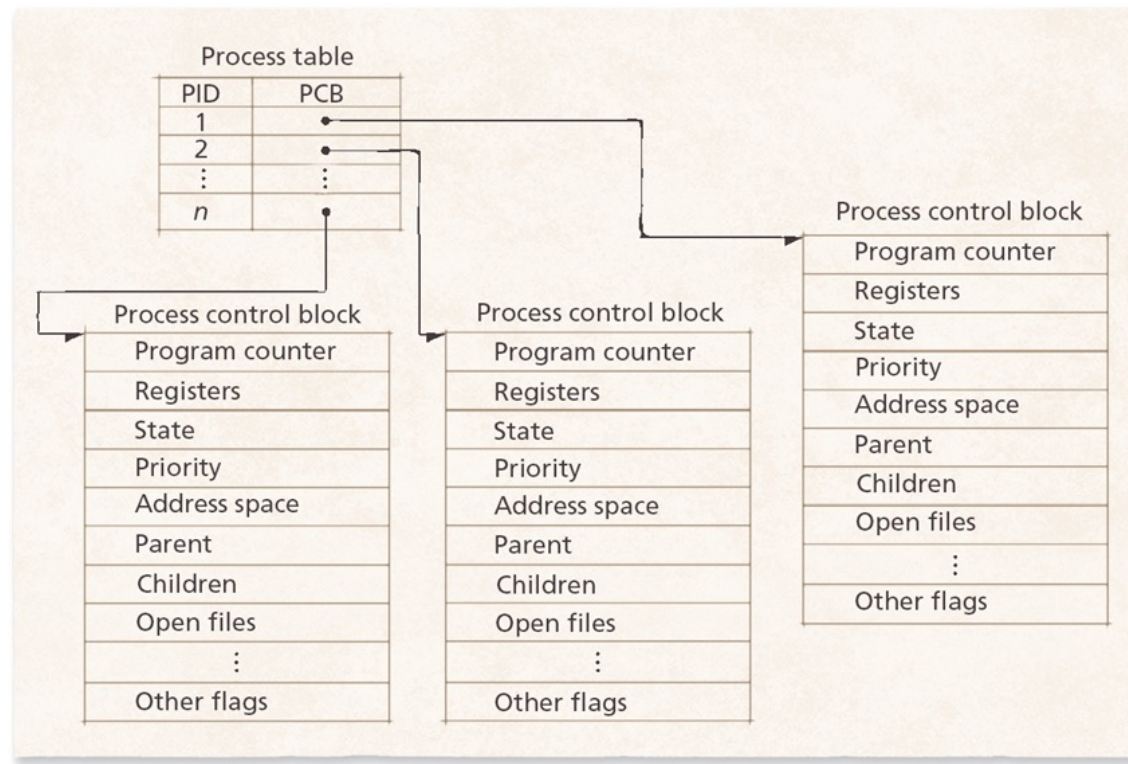
:

Example: Linux process descriptor
struct task_struct
in `/usr/src/linux/include/linux/sched.h`
at around 450 lines of statements

Process Table

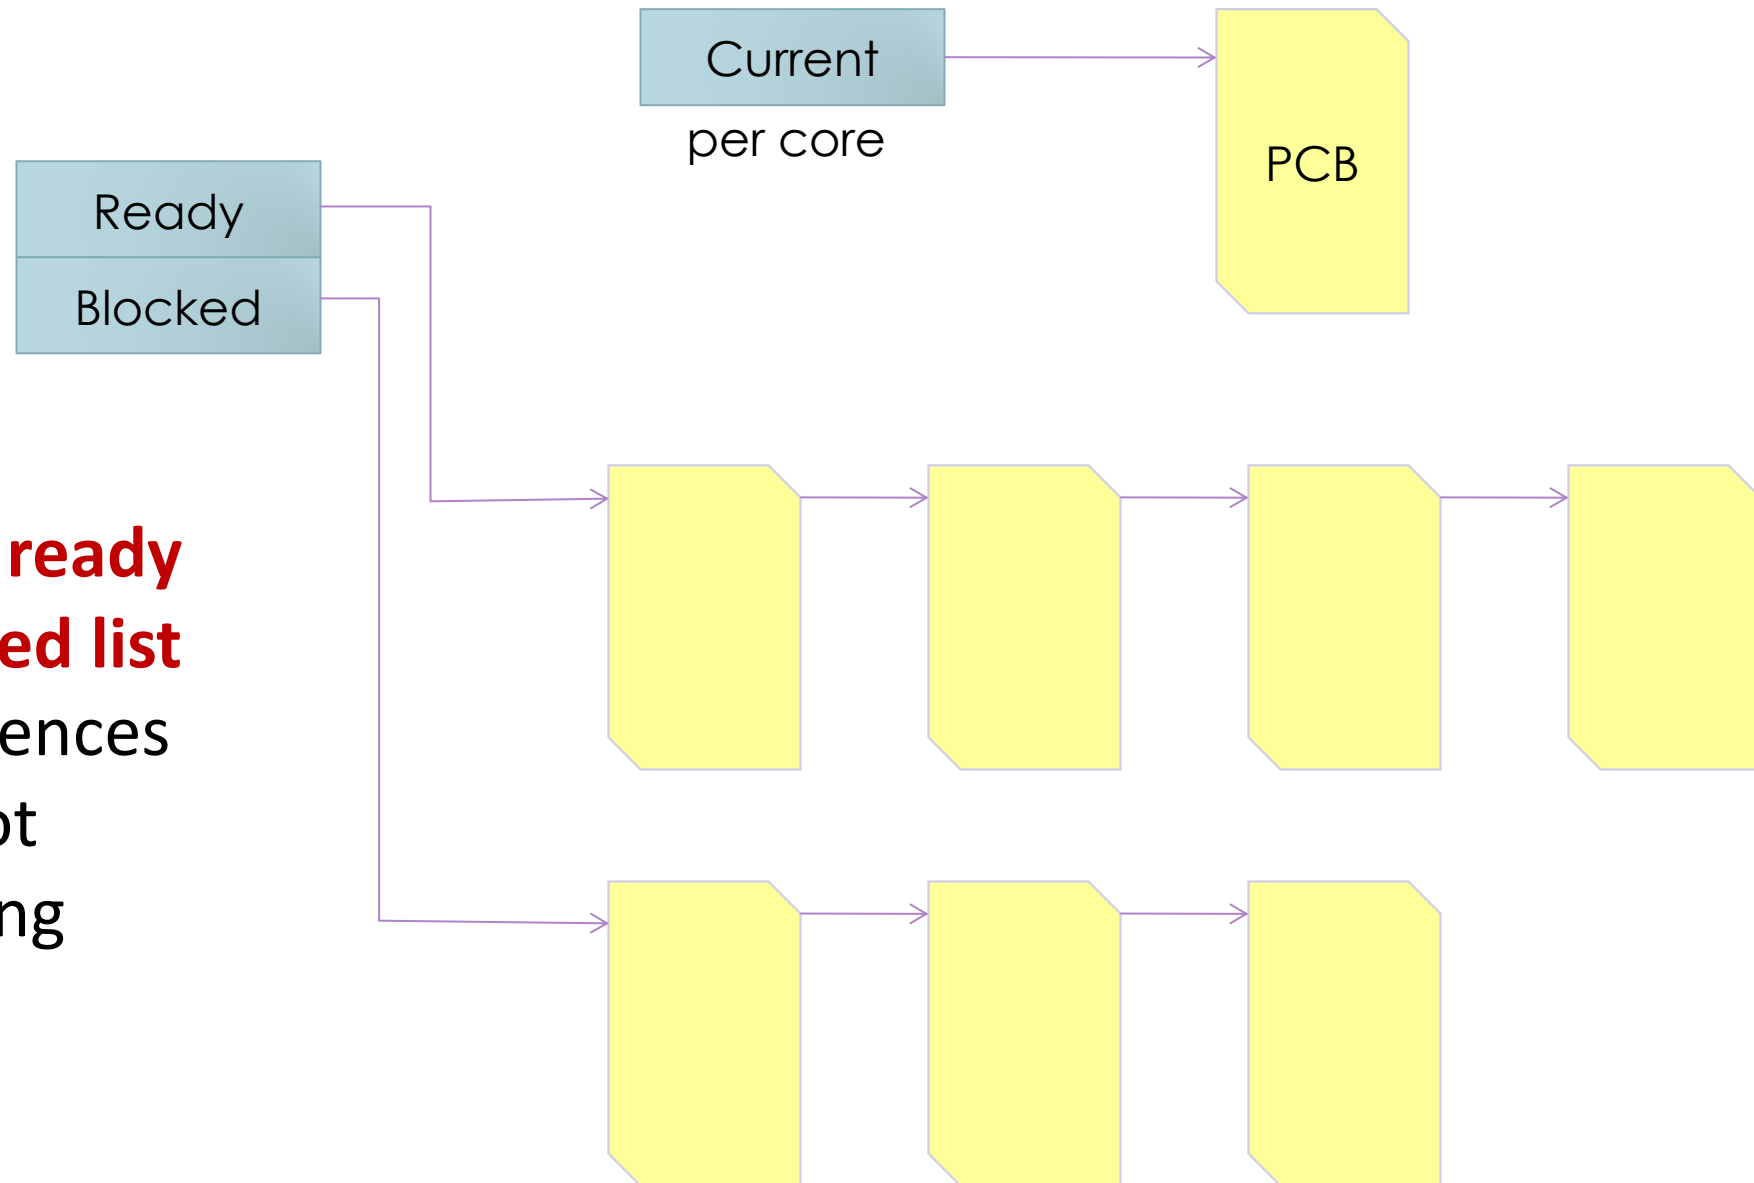
- ◉ To manage many processes, OS needs someway to **quickly access** process's PCB

- ◉ OS keeps pointers to each process's PCB in a table
 - ◉ Example - Linux "process table" is organized in a form of hashed table



- ◉ When a process is "completely" terminated, OS removes the process from the process table and frees all of the process's resources

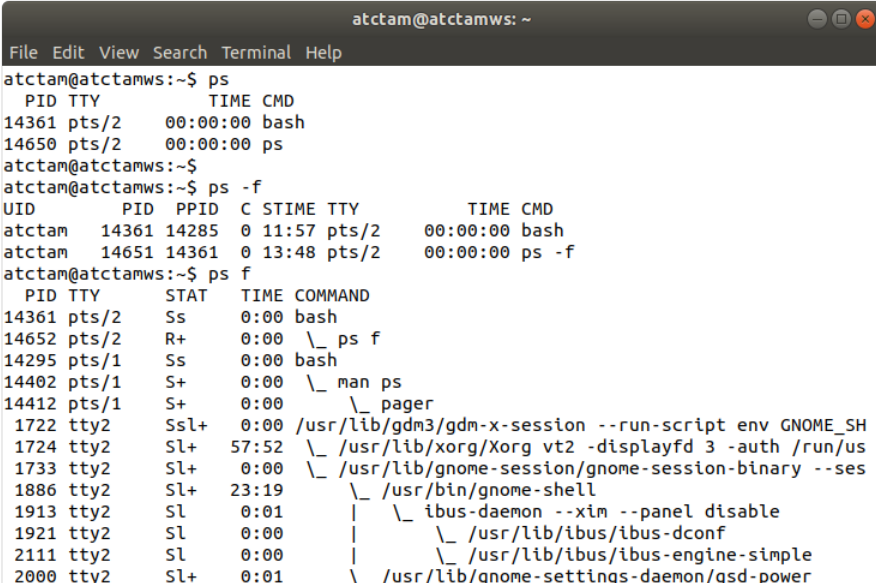
Process List Structures



OS maintains a **ready list** and a **blocked list** that store references to processes not currently running

ps - Show Processes information

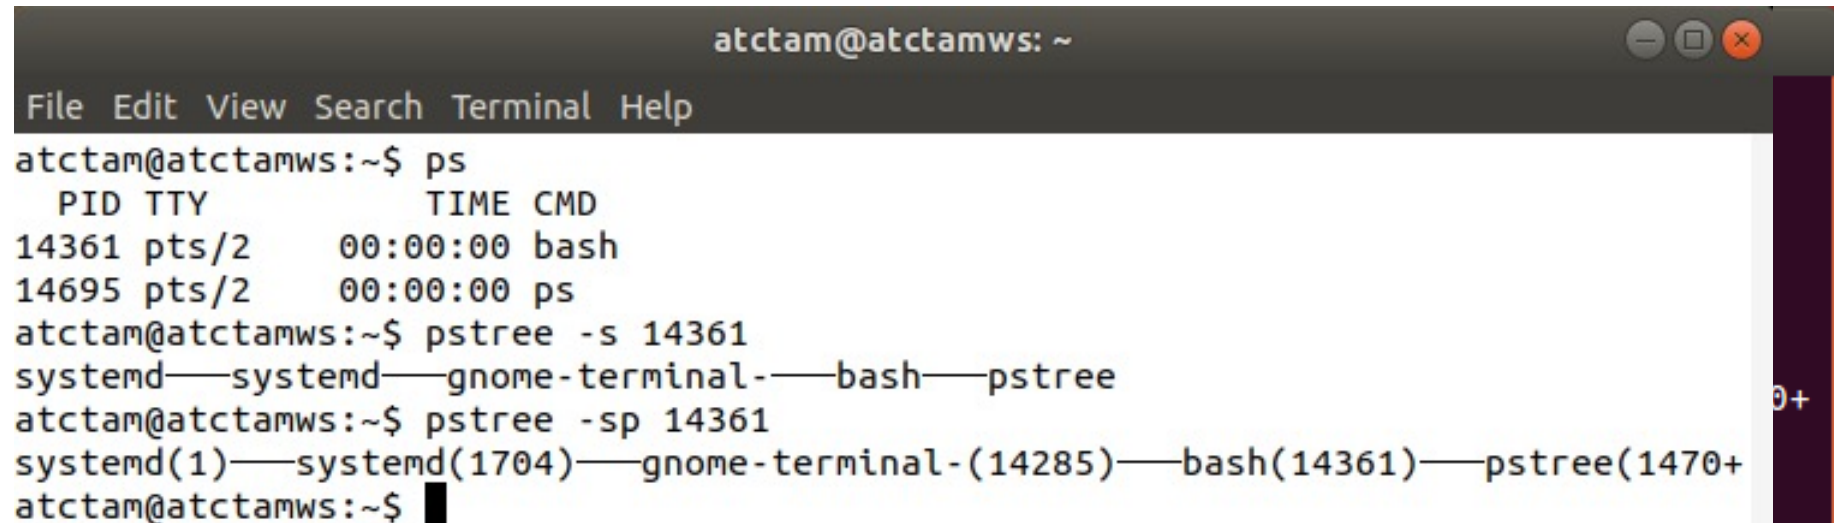
- ◉ You can list the processes' details in Linux/Mac OS X by using `ps` command
- ◉ Usage: `ps [option]`
 - ◉ When executed without any options, only processes that are associated with the current terminal are shown.
- ◉ Use man-page to learn how to use
- ◉ Some Useful options (Linux)
 - ◉ `-e`: Select all processes
 - ◉ `-f`: Show in full format
 - ◉ `w`: Wide output
 - ◉ `f`: ASCII-art process hierarchy (forest)



```
atctam@atctamws: ~  
File Edit View Search Terminal Help  
atctam@atctamws:~$ ps  
  PID TTY          TIME CMD  
14361 pts/2    00:00:00 bash  
14650 pts/2    00:00:00 ps  
atctam@atctamws:~$  
atctam@atctamws:~$ ps -f  
UID          PID  PPID  C  STIME TTY          TIME CMD  
atctam      14361 14285  0  11:57 pts/2    00:00:00 bash  
atctam      14651 14361  0  13:48 pts/2    00:00:00 ps -f  
atctam@atctamws:~$ ps f  
  PID TTY          STAT TIME COMMAND  
14361 pts/2      Ss   0:00 bash  
14652 pts/2      R+   0:00 \_ ps f  
14295 pts/1      Ss   0:00 bash  
14402 pts/1      S+   0:00 \_ man ps  
14412 pts/1      S+   0:00 \_ pager  
1722 tty2      Ssl+ 0:00 /usr/lib/gdm3/gdm-x-session --run-script env GNOME_SH  
1724 tty2      Sl+  57:52 \_ /usr/lib/xorg/Xorg vt2 -displayfd 3 -auth /run/us  
1733 tty2      Sl+  0:00 \_ /usr/lib/gnome-session/gnome-session-binary --ses  
1886 tty2      Sl+  23:19 \_ /usr/bin/gnome-shell  
1913 tty2      Sl   0:01 | \_ ibus-daemon --xim --panel disable  
1921 tty2      Sl   0:00 | \_ /usr/lib/ibus/ibus-dconf  
2111 tty2      Sl   0:00 | \_ /usr/lib/ibus/ibus-engine-simple  
2000 tty2      Sl+  0:01 \_ /usr/lib/gnome-settings-daemon/gsd-power
```

pstree - Show the relations between processes

- ◉ A Linux system program similar to *ps*
 - ◉ Processes are organized in hierarchy
 - ◉ As every process has a parent process, their relationship can be viewed by using *pstree*
- ◉ Usage: *pstree* [option]
 - ◉ Use man-page to learn how to use
 - ◉ Example:

A terminal window titled 'atctam@atctamws: ~' with a menu bar (File, Edit, View, Search, Terminal, Help). The terminal shows the output of 'ps' and 'pstree' commands. The 'ps' command shows two processes: 'bash' (PID 14361) and 'ps' (PID 14695). The 'pstree -s 14361' command shows a tree where 'systemd' is the parent of 'systemd', which is the parent of 'gnome-terminal', which is the parent of 'bash', which is the parent of 'pstree'. The 'pstree -sp 14361' command shows the same tree with PIDs in parentheses: 'systemd(1)---systemd(1704)---gnome-terminal-(14285)---bash(14361)---pstree(1470+...'.

```
atctam@atctamws: ~
File Edit View Search Terminal Help
atctam@atctamws:~$ ps
  PID TTY          TIME CMD
 14361 pts/2        00:00:00 bash
 14695 pts/2        00:00:00 ps
atctam@atctamws:~$ pstree -s 14361
systemd---systemd---gnome-terminal---bash---pstree
atctam@atctamws:~$ pstree -sp 14361
systemd(1)---systemd(1704)---gnome-terminal-(14285)---bash(14361)---pstree(1470+
atctam@atctamws:~$
```


Operations on Processes

- ◉ Operating systems provide fundamental services to processes including:
 - ◉ Creating processes
 - ◉ Destroying processes
 - ◉ Suspending processes
 - ◉ Resuming processes
 - ◉ Changing process's priority (for scheduling)
 - ◉ Waiting for a process (parent process waits for the child process)
 - ◉ Check process's status
 - ◉ Interprocess communication (IPC)
 - :

How to create a new process?

- ⦿ A process spawns a new process
 - ⦿ The process that creates a new process is now called the **parent** process
 - ⦿ The newly created process is called the **child** process
 - ⦿ It also can create other processes, thus forming a tree of processes
- ⦿ When the parent process is destroyed, modern operating systems typically responds in this way
 - ⦿ Keeps the child processes and allows them to proceed independently of the terminated parent process

Creating process

- ⦿ Actions taken
 - ⦿ Assign a unique process ID
 - ⦿ Allocate memory for the process
 - ⦿ Space for PCB must be allocated
- ⦿ Initialize the process control block
 - ⦿ Save the process ID, parent ID
 - ⦿ Set program counter and stack pointer to appropriate values
- ⦿ Set links so it is in the appropriate queue
- ⦿ Create other data structures
 - ⦿ Memory, files, accounting
- ⦿ Set the process state to Ready and put it to the Ready queue

How to Create a Process

- ◉ Unix

- ◉ fork() system function
 - ◉ A system call that **creates** a new process by **duplicating** the calling process
- ◉ exec () family of functions
 - ◉ OS replaces the current program image with a new program image

- ◉ Windows API

- ◉ CreateProcess() function
 - ◉ Creates a new process and its primary thread and loads program for execution
 - ◉ [http://msdn.microsoft.com/en-us/library/ms682512\(VS.85\).aspx](http://msdn.microsoft.com/en-us/library/ms682512(VS.85).aspx)
- ◉ CreateProcess() similar to fork() + exec ()

- ◉ Question: Why such a design of separate fork() and exec()?

Process Termination

- ◉ Process executes last statement and asks the operating system to delete itself
 - ◉ `exit()` or `return` from `main()`
- ◉ A process may terminate **involuntarily**
 - ◉ Parent may terminate execution of children processes (by sending a termination **signal**)
 - ◉ A number of **error and fault** conditions can lead to termination of process
- ◉ Return termination status from child to parent
 - ◉ Parent can obtain this info by calling `wait()`, `waitid()` or `waitpid()`
 - ◉ Process' resources are **de-allocated** by OS **afterward**

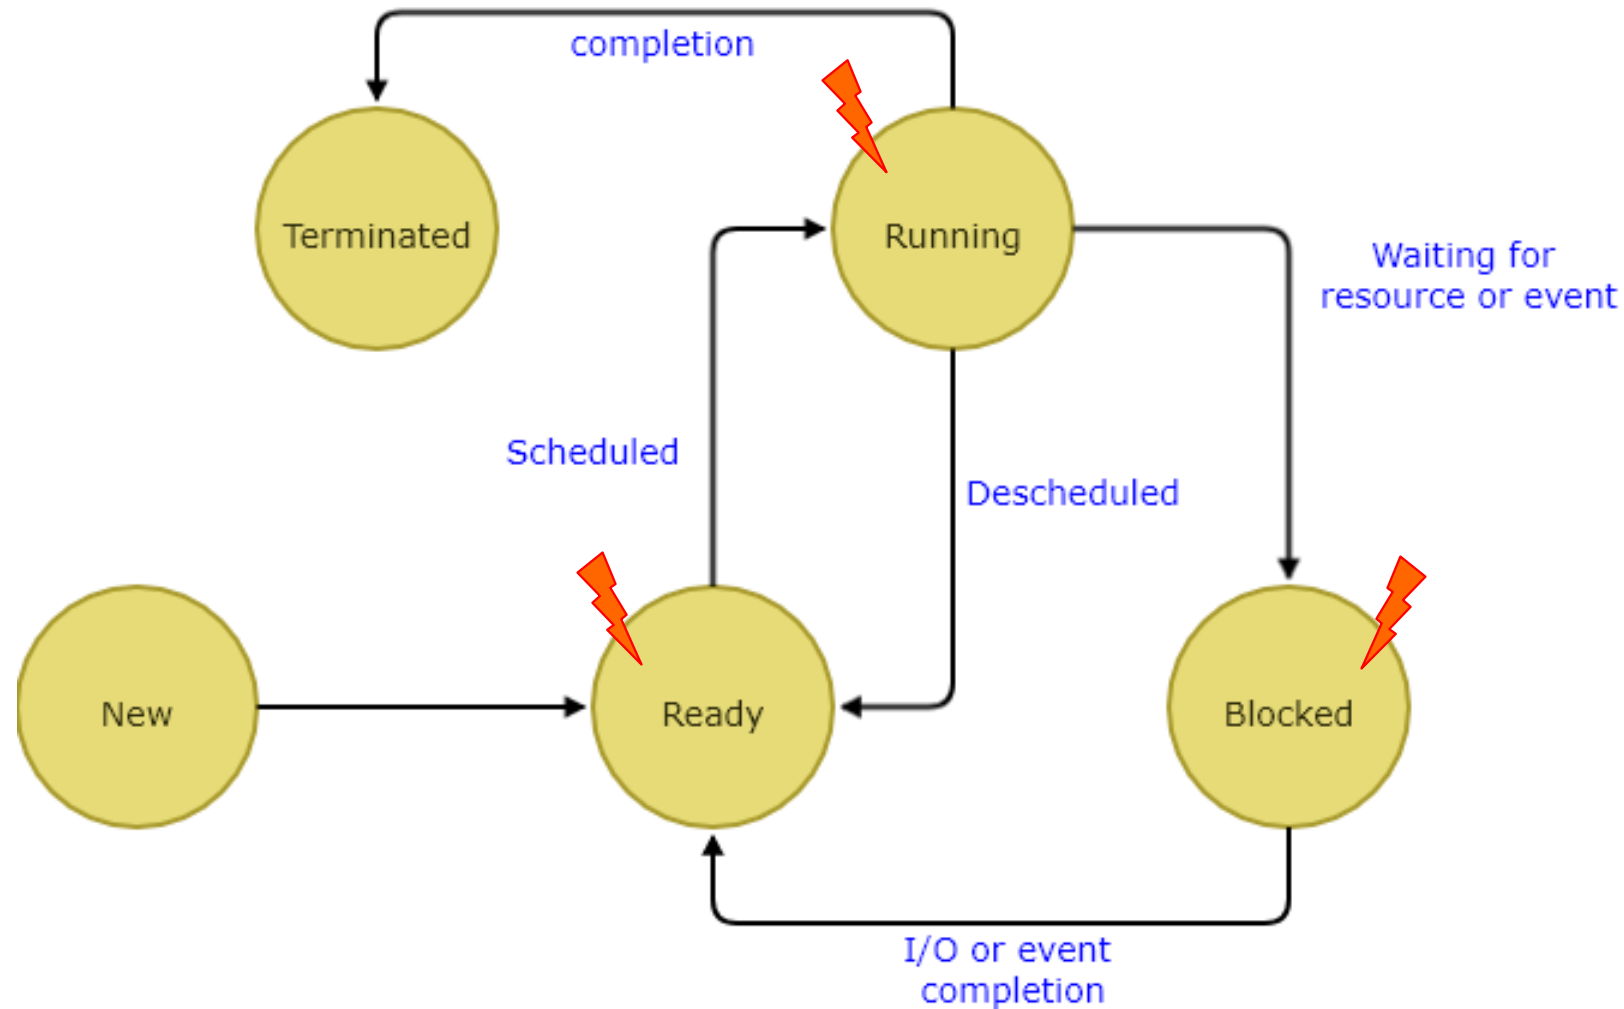
Zombie Process (UNIX)

- ⦿ A process is in terminated state
 - ⦿ When a process exits, OS still **keeps** the PCB of the process, so that the **parent can get** the information later
 - ⦿ exit status
 - ⦿ resource usage
 - ⦿ **Before** the parent process collects the information by calling **wait()** or **waitpid()**, the deceased process is kept in the Terminated state, and we called it as the “zombie” process

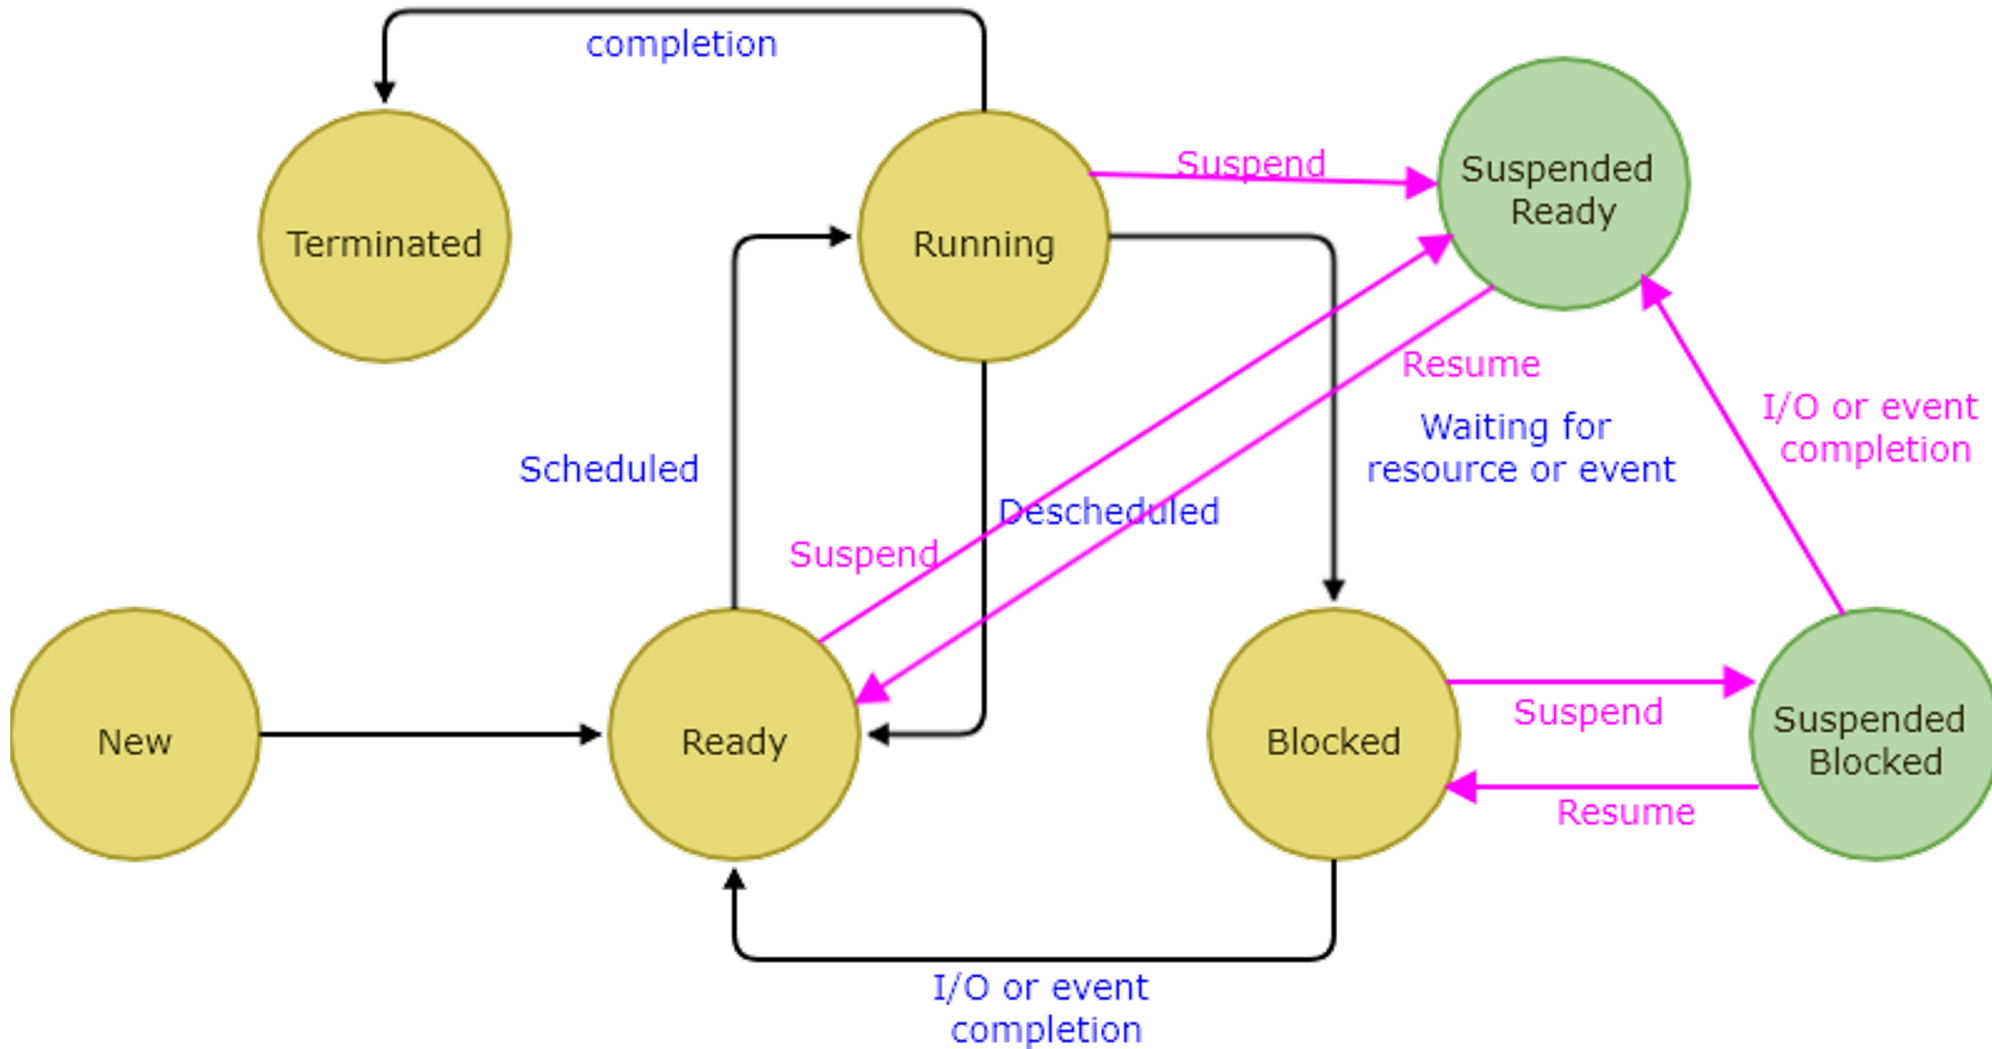
Suspending a Process

- ◉ Temporarily deactivate the process, such that it is not being considered for processor scheduling
- ◉ Why doing so:
 - ◉ Upon user request
 - ◉ Request by the parent process
 - ◉ OS may decide to suspend a **blocked** process so as to free up the memory for another ready process
- ◉ A suspended process **must be** resumed by another process
- ◉ Difference between suspension and blocked
 - ◉ **blocking** is triggered by **internal activity** of the process, while **suspend** is coming from **external**

Suspending a Process



Suspending a Process

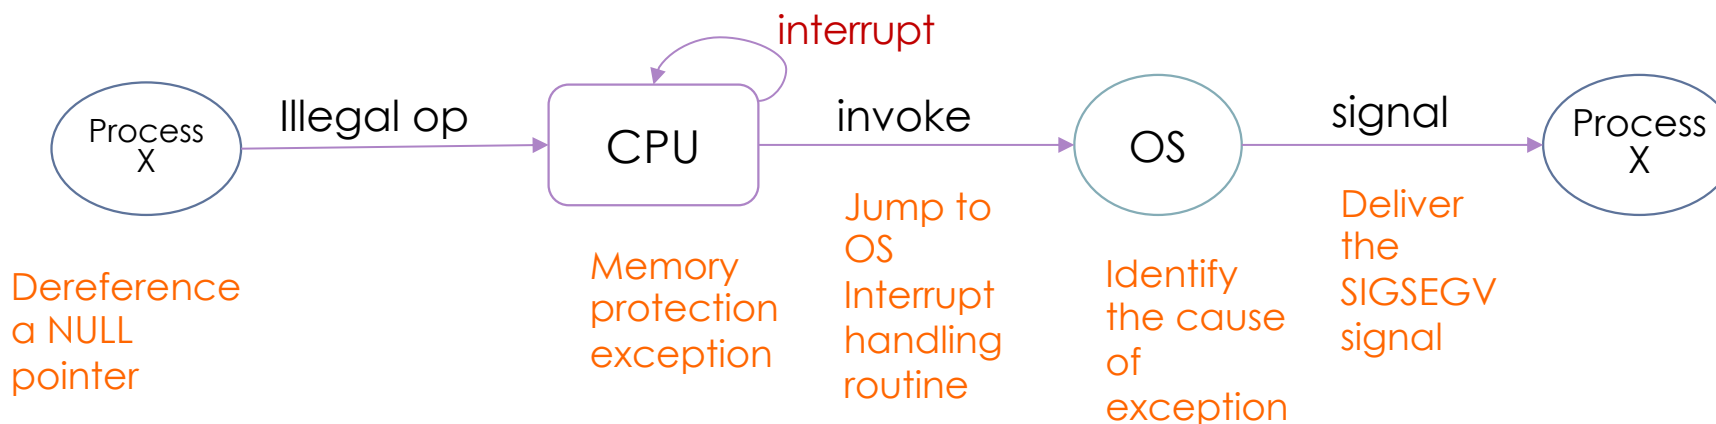


Signals

- ◉ Signals are used in UNIX systems to **notify** a process that a **particular event** has occurred
 - ◉ Sometimes being referred as “**Software Interrupt**”
- ◉ Basically implemented as system calls (kill(), signal(), sigaction(), raise(), pause(), sigsuspend(), etc.)
- ◉ Each signal is represented by a value/symbolic name
 - ◉ SIGINT: value = 2, generated when Ctrl-C is pressed
 - ◉ SIGCHLD: value = 17, generated when child process finishes execution or is terminated


Signals (2)

- A signal is generated by one **software entity** (in the occurrence of an event) to a **target software entity**
 - **Synchronous signal** – is triggered by the **current instruction** of the current running process itself and **is delivered** to **that** process by the OS immediately
 - e.g., illegal memory access, division by zero



Signals (3)

- ⊙ **Asynchronous signals** – are generated by external events / activities, which are **not triggered by the current activity** / action of the target process at the time of receiving the signal
 - ⊙ i.e., arrive at unpredictable times during execution of the program
 - ⊙ e.g., by the timer alarm
 - ⊙ e.g., parent process using kill() system call to kill the child process



Signals (4)

- ◉ A process can decide whether it wants to **catch**, **ignore** or **mask** a signal
 - ◉ **Catching** a signal involves specifying a routine (**signal handler**) in advance so that the OS will invoke that handler when the process receives that signal
 - ◉ the **signal()** or **sigaction()** system calls can be used by the program to specify the signal handler routine to the OS
 - ◉ **Catching** – Using OS's default action to handle the signal
 - ◉ **Ignore** – Inform OS that it does not want to handle that signal
 - ◉ **Masking** a signal is to instruct the OS not to deliver signals of that type until the process clears the signal mask
- ◉ SIGKILL and SIGSTOP cannot be caught, blocked or ignored

Signals (5)

- ◉ How OS determines what a process will respond to a particular signal
 - ◉ A process's PCB contains a pointer to **a vector of signal handlers** (logically order by the signal number)
 - ◉ Each entry corresponds to the handler function for that entry (signal)
- ◉ A child process inherits the setting from its parent
- ◉ However, if use `exec...()` function to load a new program image, any signals that have the custom-made handlers **will be reset** to default setting

Summary

- ◉ To manage and control a running process (application), OS needs some mechanisms to keep track on the current status of the process
 - ◉ Various data structures are needed
 - ◉ Process control block (PCB)
 - ◉ Process table and lists
- ◉ OS provides a set of operations for us to work with processes
- ◉ Signals were introduced in Unix systems to allow interactions between User Mode processes; the kernel also uses them to notify processes of system events.